














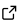

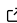
1 Jobflow: Computational Workflows Made Simple

2 **Andrew S. Rosen** ^{1,2}, **Max Gallant** ^{1,2}, **Janine George** ^{3,4}, **Janosh**
3 **Riebesell** ^{2,5}, **Hrushikesh Sahasrabudde** ^{1,6}, **Jimmy-Xuan Shen** ⁷,
4 **Mingjian Wen** ⁸, **Matthew L. Evans** ^{9,10}, **Guido Petretto**⁹, **David**
5 **Waroquiers** ^{9,10}, **Gian-Marco Rignanese** ^{9,10,11}, **Kristin A. Persson** ^{1,2,12},
6 **Anubhav Jain** ⁶, and **Alex M. Ganose** ¹³

7 **1** Department of Materials Science and Engineering, University of California, Berkeley, Berkeley, CA,
8 USA **2** Materials Science Division, Lawrence Berkeley National Laboratory, Berkeley, CA, USA **3** Federal
9 Institute for Materials Research and Testing, Department Materials Chemistry, Berlin, Germany **4**
10 Friedrich Schiller University Jena, Institute of Condensed Matter Theory and Solid-State Optics, Jena,
11 Germany **5** Department of Physics, University of Cambridge, Cambridge, UK **6** Energy Storage and
12 Distributed Resources Division, Lawrence Berkeley National Laboratory, Berkeley, CA, USA **7** Materials
13 Science Division, Lawrence Livermore National Laboratory, Livermore, CA, USA **8** William A. Brookshire
14 Department of Chemical and Biomolecular Engineering, University of Houston, Houston, TX, USA **9**
15 Matgenix SRL, rue Armand Bury 185, 6534 Gozée, Belgium **10** Institut de la Matière Condensée et des
16 Nanosciences, Université catholique de Louvain, Chemin des Étoiles 8, Louvain-la-Neuve 1348, Belgium
17 **11** School of Materials Science and Engineering, Northwestern Polytechnical University, No. 127 Youyi
18 West Road, Xi'an 710072 Shaanxi, PR China **12** Molecular Foundry, Lawrence Berkeley National
19 Laboratory, Berkeley, CA, USA **13** Department of Chemistry, Imperial College London, London, UK

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: [Open Journals](#) 

Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

License

Authors of papers retain copyright
and release the work under a
Creative Commons Attribution 4.0
International License ([CC BY 4.0](#)).

20 Summary

21 We present Jobflow, a domain-agnostic Python package for writing computational workflows
22 tailored for high-throughput computing applications. With its simple decorator-based approach,
23 functions and class methods can be transformed into compute jobs that can be stitched
24 together into complex workflows. Jobflow fully supports dynamic workflows where the full
25 acyclic graph of compute jobs is not known until runtime, such as compute jobs that launch
26 other jobs based on the results of previous steps in the workflow. The results of all Jobflow
27 compute jobs can be easily stored in a variety of filesystem- and cloud-based databases without
28 the data storage process being part of the underlying workflow logic itself. Jobflow has been
29 intentionally designed to be fully independent of the choice of workflow manager used to
30 dispatch the calculations on remote computing resources. At the time of writing, Jobflow
31 workflows can be executed either locally or across distributed compute environments via an
32 adapter to the FireWorks package, and Jobflow fully supports the integration of additional
33 workflow execution adapters in the future.

34 Statement of Need

35 The current era of big data and high-performance computing has emphasized the significant
36 need for robust, flexible, and scalable workflow management solutions that can be used to
37 efficiently orchestrate scientific calculations ([Ben-Nun et al., 2020](#); [Silva et al., 2023](#)). To date,
38 a wide variety of workflow systems have been developed, and it has become clear that there is
39 no one-size-fits-all solution due to the diverse needs of the computational community ([Al-Saadi](#)
40 [et al., 2021](#); [Existing Workflow Systems, n.d.](#)). While several popular software packages in
41 this space have emerged over the last decade, many of them require the user to tailor their
42 domain-specific code with the underlying workflow management framework closely in mind.
43 This can be a barrier to entry for many users and puts significant constraints on the portability

44 of the underlying workflows.


45 Here, we introduce Jobflow: a free, open-source Python library that makes it simple to
46 transform collections of functions into complex workflows that can be executed either locally
47 or across distributed computing environments. Jobflow has been intentionally designed to act
48 as middleware between the user's domain-specific routines that they wish to execute and the
49 workflow "manager" that ultimately orchestrates the calculations across different computing
50 environments. Jobflow uses a simple decorator-based syntax that is similar to that of other
51 recently developed workflow tools (Babuji et al., 2019; Covalent, 2023; Prefect, n.d.; Redun,
52 2023). This approach makes it possible to turn virtually any function into a Jobflow Job
53 instance (i.e., a discrete unit of work) with minimal changes to the underlying code itself.

54 Jobflow has grown out of a need to carry out high-throughput computational materials science
55 workflows at scale as part of the Materials Project (Jain et al., 2013). As the kinds of
56 calculations — from *ab initio* to semi-empirical to those based on machine learning — continue
57 to evolve and the resulting data streams continue to diversify, it was necessary to rethink how
58 we managed an increasingly diverse range of computational workflows. Going forward, Jobflow
59 will become the computational backbone of the Materials Project, which we hope will inspire
60 additional confidence in the readiness of Jobflow for production-quality scientific computing
61 applications.

62 Features and Implementation

63 Overview

64 As a simple demonstration, the example below shows how one can construct a simple Flow
65 (i.e., a graph of interdependent Jobs) composed of two sequential Jobs: the first Job adds
66 two numbers together ($1 + 2$), and the second Job multiplies the result by another number
67 ($3 * 3$). For simplicity, this Flow is executed locally, but it can be easily dispatched to a
68 remote computing environment via the selection of a different workflow manager with no
69 modifications to the underlying function definitions. While trivial, this example demonstrates
70 the simplicity of the Jobflow syntax and how it can be used to construct complex workflows
71 from user-defined functions. Note that each Job object is not run when instantiated; rather,
72 an OutputReference associated with the Job (presented to the user via a Universally Unique
73 Identifier, or UUID) is returned, which is resolved when the workflow is ultimately executed.



74

```
from jobflow import Flow, job, run_locally

@job
def add(a, b):
    return a + b

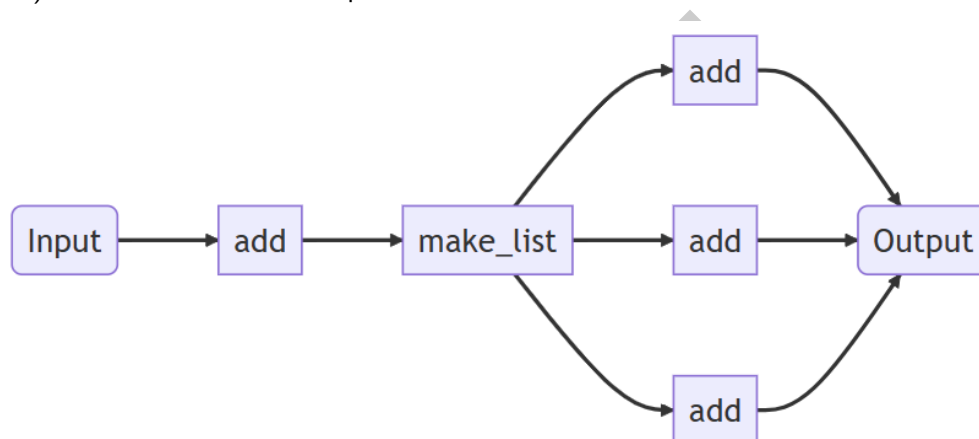
@job
def multiply(a, b):
    return a * b

job1 = add(1, 2) # 1 + 2 = 3
job2 = multiply(job1.output, 3) # 3 * 3 = 9
flow = Flow([job1, job2])

responses = run_locally(flow)
```

75 Dynamic Workflows

76 Beyond the typical acyclic graph of jobs, Jobflow fully supports dynamic workflows where the
77 precise number of jobs is unknown until runtime. This is a particularly common requirement in
78 chemistry and materials science workflows and is made possible through the use of a Response
79 object that controls the flow execution. For instance, the example below is a Flow that will add
80 two numbers ($1 + 2$), construct a list of random length containing the prior result (e.g. $[3,$
81 $3, 3]$), and then add an integer to each element of the list ($[3 + 10, 3 + 10, 3 + 10]$).
82 The `Response(replace=Flow(jobs))` syntax tells Jobflow to replace the current Job with a
83 (sub)workflow after the Job completes.



84

```

from random import randint
from jobflow import Flow, Response, job, run_locally

@job
def add(a, b):
    return a + b

@job
def make_list(val):
    return [val] * randint(2, 6)

@job
def add_distributed(vals, c):
    jobs = [add(val, c) for val in vals]
    return Response(replace=Flow(jobs))

job1 = add(1, 2) # 1 + 2 = 3
job2 = make_list(job1.output) # e.g., [3, 3, 3]
job3 = add_distributed(job2.output, 10) # [3 + 10, 3 + 10, 3 + 10]
flow = Flow([job1, job2, job3])

responses = run_locally(flow)
  
```

85 Data Management

86 Jobflow has first-class support for a variety of data stores through an interface with the magma
87 Python package ([Magma, 2023](#)). This makes it possible to easily store the results of workflows
88 in a manner that is independent of the choice of storage medium and that is entirely decoupled
89 from the workflow logic itself. Additionally, it is possible within Jobflow to specify multiple
90 types of data stores for specific Python objects (e.g., primitive types vs. large binary blobs)

91 created by a given workflow, which is often useful for storing a combination of metadata
92 (e.g., in a NoSQL database like MongoDB or file-system based store like MontyDB ([MontyDB](#),
93 [2023](#)')) and raw data (e.g., in a cloud object store like Amazon S3 or Microsoft Azure).

94 Promoting Code Reuse

95 Unlike most workflow solutions that rely on the use of functional programming, Jobflow fully
96 supports and encourages the use of compute jobs that involve class-based inheritance to reduce
97 duplication of code. While subtle, this oft-overlooked feature is particularly useful for scientific
98 workflows where very similar calculations need to be carried out but with slightly different
99 parameters or implementation details.

100 In particular, Jobflow has an abstract class called a Maker that makes it convenient to define
101 a class that can return a Job to be executed. This makes it possible to take advantage of
102 the benefits of object-oriented programming while still being able to use the straightforward
103 Jobflow decorator syntax. The support for classes also avoids the need for each workflow to
104 accept a large number of keyword arguments in order to give the user freedom to modify the
105 behavior of any constituent Job in the Flow. Instead, the class variables of the Maker can be
106 updated directly, and these changes will be reflected in the Job at runtime, as demonstrated
107 in the example below. Inheriting from the Maker class also enables updating the parameters
108 of specific Jobs in a Flow through a convenient `Flow.update_maker_kwargs(...)` function,
109 which allows for easy customization of workflows even after the Jobs have been defined.

```
from dataclasses import dataclass
from jobflow import job, Flow, Maker
from jobflow.managers.local import run_locally

@dataclass
class ExponentiateMaker(Maker):
    name: str = "Exponentiate"
    exponent: int = 2

    @job
    def make(self, a):
        return a**self.exponent

job1 = ExponentiateMaker().make(a=2) # 2**2 = 4
job2 = ExponentiateMaker(exponent=3).make(job1.output) # 4**3 = 64
flow = Flow([job1, job2])

responses = run_locally(flow)
```

110 Workflow Execution

111 One of the major benefits of Jobflow is that it decouples the details related to workflow
112 execution from the workflow definitions themselves. The simplest way to execute a workflow is
113 to run it directly on the machine where the workflow is defined using the `run_locally(...)`
114 function, as shown in the examples above. This makes it possible to quickly test even complex
115 workflows without the need to rely on a database or configuring remote resources.

116 When deploying production calculations, workflows often need to be dispatched to large
117 supercomputers through a remote execution engine. Jobflow has an interface with the
118 FireWorks package ([Jain et al., 2015](#)) via a one-line command to convert a Flow and its
119 underlying Job objects into the analogous FireWorks Workflow and Firework objects that
120 enable execution on high-performance computing machines. The logic behind the Job and
121 Flow objects are not tied to FireWorks in any direct way, such that the two packages are fully
122 decoupled.

123 Additionally, a remote mode of execution built solely around Jobflow is currently under active
124 development. With this approach, workflows can be executed across multiple “workers” (e.g.,
125 a simple computer, a supercomputer, or a cloud-based service) and managed through a
126 modern command-line interface without relying on an external workflow execution engine.
127 The forthcoming Jobflow remote mode of execution has been designed such that no inbound
128 connection from the workers to the database of jobs and results is needed, thus ensuring data
129 and network security for professional usage.

130 More generally, it is possible for users to develop custom “adapter” interfaces to their personal
131 workflow execution engine of choice. As a result, Jobflow fills a niche in the broader workflow
132 community and can help make the same workflow definition interoperable across multiple
133 workflow execution engines.

134 Testing and Documentation

135 Jobflow has been designed with robustness in mind. The Jobflow codebase has 100% test
136 coverage at the time of writing and is fully documented. The detailed testing suite, along
137 with continuous integration pipelines on GitHub, makes it easy for users to write their own
138 workflows with confidence that they will continue to work as expected for the foreseeable
139 future. Furthermore, the ability to run Jobflow Flow objects locally makes it simple to write
140 unit tests when designing a new Python package built around Jobflow without the need for
141 complex monkey-patching or spinning up a test server.

142 Usage To-Date

143 While domain-agnostic, Jobflow has been used in several materials science Python packages to
144 date, including but not limited to:

- 145 ■ Atomate2 ([Atomate2, 2023](#)), Quacc ([Quacc – the Quantum Accelerator, 2023](#)): Libraries
146 of computational chemistry and materials science workflows.
- 147 ■ NanoParticleTools ([NanoParticleTools, 2023](#)): Workflows for Monte Carlo simulations of
148 nanoparticles.
- 149 ■ Reaction Network ([McDermott et al., 2021](#); [Reaction Network, 2023](#)): Workflows for
150 constructing and analyzing inorganic chemical reaction networks.
- 151 ■ WFacer ([WFacer, 2023](#)): Workflows for modeling the statistical thermodynamics of
152 solids via automated cluster expansion.

153 Additional Details

154 Naturally, the summary presented in this article constitutes only a small subset of the features
155 that Jobflow has to offer. For additional details along with helpful tutorials ranging from basic
156 applications to examples specifically targeting the computational materials science community,
157 we refer the reader to the [Jobflow documentation](#). Suggestions, contributions, and bug reports
158 are always welcome.

159 Acknowledgements

160 This work was primarily funded and intellectually led by the Materials Project, which is funded
161 by the U.S. Department of Energy, Office of Science, Office of Basic Energy Sciences, Materials
162 Sciences and Engineering Division, under Contract no. DE-AC02-05-CH11231: Materials
163 Project program KC23MP. A.S.R. acknowledges support via a Miller Research Fellowship
164 from the Miller Institute for Basic Research in Science, University of California, Berkeley. J.G
165 would like to acknowledge the Gauss Centre for Supercomputing e.V. (www.gauss-centre.eu)
166 for funding workflow-related developments by providing generous computing time on the

167 GCS Supercomputer SuperMUC-NG at Leibniz Supercomputing Centre (www.lrz.de) (Project
168 pn73da). J.R. acknowledges support from the German Academic Scholarship Foundation
169 (Studienstiftung). M.L.E. thanks the BEWARE scheme of the Wallonia-Brussels Federation
170 for funding under the European Commission's Marie Curie-Skłodowska Action (COFUND
171 847587). G.P. and D.W. acknowledge Umicore for the financial support in developing the
172 remote execution mode of jobflow. D.W. and G.M.R. acknowledge funding from the European
173 Union's Horizon 2020 research and innovation program under the grant agreement No 951786
174 (NOMAD CoE). A.M.G. is supported by EPSRC Fellowship EP/T033231/1.

175 References

- 176 Al-Saadi, A., Ahn, D. H., Babuji, Y., Chard, K., Corbett, J., Hategan, M., Herbein, S.,
177 Jha, S., Laney, D., Merzky, A., & others. (2021). Exaworks: Workflows for exascale.
178 *2021 IEEE Workshop on Workflows in Support of Large-Scale Science (WORKS)*, 50–57.
179 <https://doi.org/10.1109/works54523.2021.00012>
- 180 *Atomate2*. (2023). <https://github.com/materialsproject/atomate2>
- 181 Babuji, Y., Woodard, A., Li, Z., Katz, D. S., Clifford, B., Kumar, R., Laciniski, L., Chard, R.,
182 Wozniak, J. M., Foster, I., & others. (2019). Parsl: Pervasive parallel programming in
183 python. *Proceedings of the 28th International Symposium on High-Performance Parallel
184 and Distributed Computing*, 25–36.
- 185 Ben-Nun, T., Gamblin, T., Hollman, D. S., Krishnan, H., & Newburn, C. J. (2020). Workflows
186 are the new applications: Challenges in performance, portability, and productivity. *2020
187 IEEE/ACM International Workshop on Performance, Portability and Productivity in HPC
188 (P3HPC)*, 57–69. <https://doi.org/10.1109/p3hpc51967.2020.00011>
- 189 *Covalent*. (2023). <https://doi.org/10.5281/zenodo.5903364>
- 190 *Existing workflow systems*. (n.d.). <https://s.apache.org/existing-workflow-systems>
- 191 Jain, A., Ong, S. P., Chen, W., Medasani, B., Qu, X., Kocher, M., Brafman, M., Petretto, G.,
192 Rignanesi, G.-M., Hautier, G., & others. (2015). FireWorks: A dynamic workflow system
193 designed for high-throughput applications. *Concurrency and Computation: Practice and
194 Experience*, 27(17), 5037–5059. <https://doi.org/10.1002/cpe.3505>
- 195 Jain, A., Ong, S. P., Hautier, G., Chen, W., Richards, W. D., Dacek, S., Cholia, S., Gunter,
196 D., Skinner, D., Ceder, G., & others. (2013). Commentary: The materials project: A
197 materials genome approach to accelerating materials innovation. *APL Materials*, 1(1).
198 <https://doi.org/10.1063/1.4812323>
- 199 *Magma*. (2023). <https://github.com/materialsproject/magma>
- 200 McDermott, M. J., Dwaraknath, S. S., & Persson, K. A. (2021). A graph-based network for
201 predicting chemical reaction pathways in solid-state materials synthesis. *Nature Communi-
202 cations*, 12(1), 3097. <https://doi.org/10.1038/s41467-021-23339-x>
- 203 *MontyDB*. (2023). <https://github.com/davidlatwe/montydb>
- 204 *NanoParticleTools*. (2023). <https://github.com/BlauGroup/NanoParticleTools>
- 205 *Prefect*. (n.d.). <https://github.com/PrefectHQ/prefect>
- 206 *Quacc – the quantum accelerator*. (2023). <https://doi.org/10.5281/zenodo.7720998>
- 207 *Reaction network*. (2023). <https://github.com/materialsproject/reaction-network>
- 208 *Redun*. (2023). <https://github.com/insitro/redun>

- 209 Silva, R. F. da, Badia, R. M., Bala, V., Bard, D., Bremer, P.-T., Buckley, I., Caino-Lores, S.,
210 Chard, K., Goble, C., Jha, S., & others. (2023). Workflows community summit 2022: A
211 roadmap revolution. *arXiv Preprint arXiv:2304.00019*.
- 212 *WFacer*. (2023). <https://github.com/CederGroupHub/WFacer>

DRAFT